KnowledgeOS Test Document — RAG Guide

Purpose

This PDF is a test input for the KnowledgeOS document ingestion pipeline.

Ingestion Process

The system stores the original document in MinIO, parses the document, extracts text, splits the text into chunks, generates embeddings, and stores the resulting vectors in PGVector.

Retrieval

When a user asks a question, KnowledgeOS retrieves relevant document chunks and supplies them to the language model so the answer can be grounded in organizational knowledge.